

Human Activity Recognition with Smartphones

Ensemble Learning Project. (By: - Shivesh (Ethan) Sahu)

Problem Definition

The goal of this project is to automatically classify **human activities** (e.g. walking, sitting, standing) based on smartphone sensor measurements, using classical machine learning and ensemble methods. The focus is not only on building a high-performing model, but also on:

- Comparing a range of linear, distance-based, and tree-based classifiers.
 - Applying **hyperparameter tuning** and **ensemble methods** (voting, stacking, bagging, boosting).
 - Understanding where the model still makes mistakes (especially between **SITTING** and **STANDING**).
 - Packaging the final pipeline into a reusable artifact (`best_har_model.joblib`) with an easy inference helper.
-

Dataset Overview

- **Dataset:** Human Activity Recognition with Smartphones (Kaggle) – `train.csv`
 - **Shape:** 7352×563
 - **561** numeric feature columns (engineered time/frequency domain features from accelerometer + gyroscope).
 - **1** subject identifier: subject (integer ID).
 - **1** target column: Activity (string label).
 - **Target classes** and counts (in the full dataset):
 - LAYING – 1407 samples
 - STANDING – 1374 samples
 - SITTING – 1286 samples
 - WALKING – 1226 samples
 - WALKING_UPSTAIRS – 1073 samples
 - WALKING_DOWNSTAIRS – 986 samples
 - **Data quality:**
 - All 561 feature columns are `float64`.
 - subject is `int64`.
 - Activity is `object` (categorical labels).
 - **No missing values:** total missing count = 0 across the entire dataset.
-

Exploratory Data Analysis (EDA)

3.1 Basic statistics

For the first 10 features (e.g. tBodyAcc-mean()-X, tBodyAcc-std()-X, ...), summary statistics indicate:

- Features are already scaled into a range around -1, 1 with mean values typically between ~ -0.7 and 0.3.
- Standard deviations vary but are moderate, confirming that features are numerically well-behaved.

3.2 Class distribution

A count plot of Activity shows:

- Slight class imbalance (e.g. more **LAYING** and **STANDING** than **WALKING_DOWNSTAIRS**), but all classes still have >900 samples.
- No extreme minority classes, so standard cross-entropy / macro-F1 evaluation is appropriate.

3.3 Subject distribution

- subject counts range from ~ 281 to 366 per user (first ten IDs inspected).
- This confirms that activities are captured across multiple individuals, which reduces the risk of the model overfitting to a single subject.

3.4 Feature behavior and correlations

For a small subset of features:

`["tBodyAcc-mean()-X", "tBodyAcc-mean()-Y", "tBodyAcc-std()-X", "tBodyAcc-std()-Y"]`

- **Histograms:**
 - Means (tBodyAcc-mean()) show fairly concentrated distributions, with some skew.
 - Standard deviations vary more widely, reflecting differences in movement intensity.
- **Boxplots:**

- Some potential outliers appear, but they are plausible sensor readings rather than obvious errors.
 - **Correlation heatmap** (subset only):
 - High positive correlation between tBodyAcc-std()-X and tBodyAcc-std()-Y (~0.93).
 - Weak correlations between mean and std features across axes (near 0), indicating they provide complementary information.
-

Data Preparation

4.1 Feature/target split

- **Features (X)**: all columns except Activity and subject.
 - 561 numeric features kept.
 - subject is deliberately removed to avoid ID-leakage (we want a model that generalizes across unseen people).
- **Target (y)**: Activity.

4.2 Train–test split

- `train_test_split` with:
 - `test_size=0.2`
 - `random_state=42`
 - `stratify=y` to preserve class proportions.

Result:

- `X_train`: (5881, 561)
- `X_test`: (1471, 561)

4.3 Scaling & pipelines

Because many algorithms are sensitive to feature scaling, a **ColumnTransformer + Pipeline** pattern was used:

- `scaler = ColumnTransformer([("num", StandardScaler(), num_features)], remainder="drop")`

- Two helper constructors:
 - *def scaled_pipeline(estimator):*
 - *return Pipeline([("scaler", scaler), ("estimator", estimator)])*
 -
 - *def passthrough_pipeline(estimator):*
 - *return Pipeline([("clf", estimator)])*
 - **Scaled** models:
 - linear and distance-based (Perceptron, RidgeClassifier, LogisticRegression, SVM, KNN).
 - **Unscaled / passthrough:**
 - tree-based models and ensembles (DecisionTree, RandomForest, bagging/boosting).
-

Baseline Models

The following baseline models were trained on the training set and evaluated on the test set using **accuracy** and **macro-averaged F1**:

Model	Accuracy	F1 (macro)
SVM (RBF)	0.9898	0.9906
RandomForest	0.9878	0.9879
Logistic Regression	0.9857	0.9868
RidgeClassifier (Adaline-like)	0.9810	0.9822
Perceptron	0.9810	0.9818
KNN (k = 15, distance)	0.9565	0.9578
DecisionTree	0.9388	0.9380

Observations:

- All linear models perform surprisingly well once features are standardized, suggesting that the engineered HAR features are highly informative and almost linearly separable.

- SVM with RBF kernel already reaches **~0.99 accuracy**, indicating minimal remaining headroom.

A confusion matrix for the **baseline SVM** shows an almost perfect diagonal; the most noticeable residual confusion is between **SITTING** and **STANDING** (some SITTING predicted as STANDING and vice versa), which is intuitively reasonable given their similar motion patterns.

Hyperparameter Tuning

To push baseline models further, **RandomizedSearchCV** was applied with stratified K-fold CV:

- `FAST_SEARCH = True` \rightarrow `cv_splits = 3`, modest `n_iter` per model (15).
- Scoring metric: `f1_macro`.
- Models tuned:
 - SVM (RBF)
 - RandomForest
 - Logistic Regression
 - KNN

Tuned parameter highlights

- **SVM (RBF)**
 - Best: $C \approx 31.43$, `gamma="auto"`
 - CV macro-F1 ≈ 0.9830
 - Test: unchanged vs baseline (~ 0.9898 acc / 0.9906 F1), i.e., already near optimal.
- **RandomForest**
 - Best: ~ 341 trees, `max_depth=30`, `max_features="log2"`, `min_samples_split=3`, `min_samples_leaf=2`.
 - CV macro-F1 ≈ 0.9757 .
 - Test: ≈ 0.9850 acc / 0.9853 F1 (slightly below the untuned 300-tree RF).
- **Logistic Regression**
 - Best: $C \approx 0.53$, `penalty="l2"`, `solver="lbfgs"`.
 - Test: ≈ 0.9857 acc / 0.9868 F1 (similar to baseline).
- **KNN**
 - Best: `n_neighbors=16`, `weights="uniform"`, `p=1`.
 - Test: ≈ 0.9708 acc / 0.9721 F1 (noticeable improvement over the original $k=15$ but still below linear/SVM models).

Conclusion: tuning delivered small gains, but SVM (RBF) and RandomForest were already very strong.

Classical Ensembles: Voting & Stacking

Using the tuned (or baseline) pipelines as building blocks, three higher-level ensembles were built:

1. **Voting (hard)** – majority vote of:
 - o Perceptron, RidgeClassifier, Logistic Regression (tuned), SVM (RBF, tuned), DecisionTree, RandomForest (tuned), KNN (tuned).
2. **Voting (soft)** – probability-average of:
 - o Logistic Regression (tuned), SVM (RBF with probability=True), RandomForest (tuned), KNN (tuned).
3. **Stacking (LR meta)** – meta-learner:
 - o Base estimators: Logistic Regression (tuned), SVM (probabilistic), RandomForest (tuned), KNN (tuned).
 - o Final estimator: Logistic Regression with max_iter=5000.
 - o stack_method="predict_proba", cross-validated meta-features with StratifiedKFold.

7.1 Ensemble performance

Ensemble	Accuracy F1 (macro)	
Voting (hard)	0.9905	0.9912
Voting (soft)	0.9918	0.9925
Stacking (LR meta)	0.9946	0.9950

The **Stacking** ensemble clearly outperforms all individual models and both voting schemes.

7.2 Confusion matrix & error analysis (Stacking)

The confusion matrix for **Stacking (LR meta)** is extremely close to a perfect diagonal:

- All **walking-type activities** (WALKING, WALKING_UPSTAIRS, WALKING_DOWNSTAIRS) are classified almost perfectly.
- **LAYING** has near-perfect precision/recall.
- Residual errors are almost entirely between **SITTING** and **STANDING** (e.g., a handful of SITTING misclassified as STANDING and vice versa).

Interpretation:

- SITTING vs STANDING are biomechanically similar when captured by a waist-mounted smartphone; differences in vertical acceleration and orientation are subtle.
 - Additional sensor context (e.g., orientation sequences, seat pressure sensors) might be needed to fully disambiguate them; given the current feature set, an error rate of only a few samples is already extremely strong.
-

Bagging & Boosting Experiments

To keep tree-based ensembles clearly separated from the original model dictionary, a new `bag_boost_models` registry was built using `passthrough_pipeline`:

- **Bagging (DT)** – BaggingClassifier with DecisionTree base learners.
- **ExtraTrees** – Extremely randomized trees with 400 trees.
- **AdaBoost** – AdaBoostClassifier with shallow DecisionTree base estimators.
- **HistGradientBoosting** – Histogram-based gradient boosting classifier.

8.1 Baseline results

Model	Accuracy	F1 (macro)
HistGradientBoosting	0.9939	0.9941
ExtraTrees	0.9932	0.9937
AdaBoost	0.9864	0.9869
Bagging (DT)	0.9721	0.9722

These results are very competitive:

- **HistGradientBoosting** and **ExtraTrees** nearly match the Stacking ensemble, confirming that boosted/bagged trees are also an excellent fit for this problem.
- However, in this run the **Stacking (LR meta) model still achieves the highest macro-F1**, so it remains the global best.

8.2 Hyperparameter tuning for bagging/boosting

A second stage of RandomizedSearchCV was run with a **QUICK** mode ($n_iter \approx 12$ for most models):

- **ExtraTrees** – tuned over number of trees, max depth, max features, and min samples split/leaf.
- **Bagging (DT)** – tuned over number of estimators, bootstrap vs. pasting, and inner tree depth / split settings.
- **AdaBoost** – tuned number of estimators, learning rate, and max depth of the base tree.
- **HistGradientBoosting** – tuned over learning rate, max iterations, max leaf nodes, max depth, min samples leaf, and L2 regularization.

The tuned versions improved macro-F1 further (see notebook output for exact values), but **none were promoted as the final production model**, primarily because:

- The stacking model was already saved and integrated into the deployment artifact.
- Differences in performance among the very top models are small enough that choosing the simpler, already-packaged model is reasonable.

Final Model Selection and Artifacts

9.1 Global leaderboard (non-bagging/boosting)

From the consolidated comparison of all non-bagging/boosting models plus ensembles:

1. **Stacking (LR meta)** – 0.9946 acc / 0.9950 macro-F1
2. Voting (soft)
3. Voting (hard)
4. SVM (RBF, tuned)

5. SVM (RBF, baseline)
6. RandomForest, etc.

Thus, **Stacking (LR meta)** is selected as the **final model** for deployment.

9.2 Saved artifacts

Two key artifacts were persisted:

1. **Model**
 - o File: `best_har_model.joblib`
 - o Content: full scikit-learn pipeline for **Stacking (LR meta)**, including all preprocessing steps (scaling) and base/meta learners.
2. **Feature list**
 - o File: `feature_names.txt`
 - o One feature name per line in the exact order used during training.
 - o Ensures that incoming data for inference is correctly aligned with the model's expected columns.

Both files are stored alongside the notebook and can be version-controlled.

Inference Interface

To make real-world usage convenient, a helper function was implemented:

```
def predict_activities(new_data, model_path=MODEL_PATH,
feature_list_path=FEATURE_LIST_PATH):
```

""""

new_data: pandas DataFrame (same feature space) OR path to CSV with same columns.

Returns: numpy array of predicted activity labels.

""""

```
from joblib import load
```

```
if isinstance(new_data, str):
```

```

new_df = pd.read_csv(new_data)

else:
    new_df = new_data.copy()

# Load feature list and check columns
with open(feature_list_path, "r") as f:
    feats = [ln.strip() for ln in f if ln.strip()]

missing = [c for c in feats if c not in new_df.columns]

if missing:
    raise ValueError(
        f"Missing required columns: {missing[:10]}\n" +
        "... if len(missing) > 10 else \""
    )

# Align column order
new_df = new_df.reindex(columns=feats)

# Load model and predict
model = load(model_path)

return model.predict(new_df)

```

Key properties:

- Accepts either a **DataFrame** or a path to a **CSV file**.
- Validates that all required features are present; fails fast with a helpful message if not.
- Reorders columns to match training order before prediction.
- Returns predicted **activity labels** as a NumPy array.

This function can be imported into a separate script or API service to perform inference on new sensor data batches.

Limitations and Future Work

1. **SITTING vs STANDING confusion**
 - o Even the best model occasionally confuses these two classes, reflecting the inherent similarity from accelerometer/gyroscope data alone.
 - o Future work could explore:
 - Additional engineered features capturing postural orientation more explicitly.
 - Sequence models (e.g. LSTMs or HMMs) that consider temporal dynamics rather than single windows.
2. **Subject leakage / generalization**
 - o The model is trained and tested on the same subject pool (though stratified).
 - o A stricter evaluation would use **subject-wise cross-validation**: train on some subjects, test on entirely unseen subjects.
3. **Model size and latency**
 - o Stacking combines several heavy base learners; for deployment on a resource-constrained device, a lighter model such as HistGradientBoosting or a single tuned SVM might be preferable.
4. **Hyperparameter search budget**
 - o All searches used a “FAST_SEARCH” / “QUICK” configuration; deeper search (more `n_iter`, 5-fold CV) could extract slightly more performance at higher compute cost.

Summary

- Using the Kaggle Human Activity Recognition dataset with 561 engineered features, we built and compared a rich set of classical ML models and ensembles.
- After careful preprocessing, EDA, baseline comparison, and hyperparameter tuning, the **StackingClassifier with Logistic Regression meta-learner** emerged as the best performer with **~0.995 macro-F1** on a held-out test set.
- Bagging/boosting methods (ExtraTrees, HistGradientBoosting, AdaBoost, Bagging DT) proved highly competitive and almost matched stacking performance, further validating the predictability of the problem.
- The project concludes with a **production-ready pipeline** (`best_har_model.joblib`), a **feature specification file**, and a **robust inference helper**, making the solution easy to integrate into downstream systems or real-time HAR applications.