

```
In [1]: #=====
# ENSEMBLE LEARNING PROJECT
# Human Activity Recognition with Smartphones (Kaggle)
#=====

# ----- Runtime knobs -----
FAST_SEARCH = True # True = faster hyperparam search; set False for

# 1. Importing all the important libraries
import warnings
warnings.filterwarnings("ignore")

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, StratifiedKFold,
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer

from sklearn.linear_model import Perceptron, RidgeClassifier, Logistic
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, VotingClassifier,
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, f1_score, classification_r

from sklearn.base import clone
from scipy.stats import loguniform, randint as sp_randint
from joblib import dump
```

```
In [2]: # ----- 1.1 Load dataset -----
df = pd.read_csv("train.csv")
print("\n== Loading the Datasets ==")
print("\nDataset Shape(rows, columns): ", df.shape)
print("\nFirst 5 rows: \n", df.head(5))
print("\nColumn names: \n", df.columns.tolist())
```

== Loading the Datasets ==

Dataset Shape(rows, columns): (7352, 563)

First 5 rows:

	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-mean()-Z	tBodyAcc-s
td()-X \				
0	0.288585	-0.020294	-0.132905	-0.9
95279				
1	0.278419	-0.016411	-0.123520	-0.9
98245				

2	0.279653	-0.019467	-0.113462	-0.9
95380				
3	0.279174	-0.026201	-0.123283	-0.9
96091				
4	0.276629	-0.016570	-0.115362	-0.9
98139				
	tBodyAcc-std()-Y	tBodyAcc-std()-Z	tBodyAcc-mad()-X	tBodyAcc-mad()
-Y \				
0	-0.983111	-0.913526	-0.995112	-0.9831
85				
1	-0.975300	-0.960322	-0.998807	-0.9749
14				
2	-0.967187	-0.978944	-0.996520	-0.9636
68				
3	-0.983403	-0.990675	-0.997099	-0.9827
50				
4	-0.980817	-0.990482	-0.998321	-0.9796
72				
	tBodyAcc-mad()-Z	tBodyAcc-max()-X	...	fBodyBodyGyroJerkMag-kurtosis()
\				
0	-0.923527	-0.934724	...	-0.71
0304				
1	-0.957686	-0.943068	...	-0.86
1499				
2	-0.977469	-0.938692	...	-0.76
0104				
3	-0.989302	-0.938692	...	-0.48
2845				
4	-0.990441	-0.942469	...	-0.69
9205				
	angle(tBodyAccMean, gravity)	angle(tBodyAccJerkMean), gravityMean) \		
0	-0.112754		0.030400	
1	0.053477		-0.007435	
2	-0.118559		0.177899	
3	-0.036788		-0.012892	
4	0.123320		0.122542	
	angle(tBodyGyroMean, gravityMean)	angle(tBodyGyroJerkMean, gravityMean) \		
0	-0.464761		-0.0184	
46				
1	-0.732626		0.7035	
11				
2	0.100699		0.8085	
29				
3	0.640011		-0.4853	
66				
4	0.693578		-0.6159	
71				

```

    angle(X,gravityMean)  angle(Y,gravityMean)  angle(Z,gravityMean)  su
bject \
0          -0.841247          0.179941        -0.058627
1          -0.844788          0.180289        -0.054317
1          -0.848933          0.180637        -0.049118
1          -0.848649          0.181935        -0.047663
1          -0.847865          0.185151        -0.043892
1

Activity
0  STANDING
1  STANDING
2  STANDING
3  STANDING
4  STANDING

[5 rows x 563 columns]

```

Column names:

```

['tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-mean()-Z', 'tBody
Acc-std()-X', 'tBodyAcc-std()-Y', 'tBodyAcc-std()-Z', 'tBodyAcc-mad()-
X', 'tBodyAcc-mad()-Y', 'tBodyAcc-mad()-Z', 'tBodyAcc-max()-X', 'tBodyA
cc-max()-Y', 'tBodyAcc-max()-Z', 'tBodyAcc-min()-X', 'tBodyAcc-min()-
Y', 'tBodyAcc-min()-Z', 'tBodyAcc-sma()', 'tBodyAcc-energy()-X', 'tBody
Acc-energy()-Y', 'tBodyAcc-energy()-Z', 'tBodyAcc-iqr()-X', 'tBodyAcc-i
qqr()-Y', 'tBodyAcc-iqr()-Z', 'tBodyAcc-entropy()-X', 'tBodyAcc-entrop
y()-Y', 'tBodyAcc-entropy()-Z', 'tBodyAcc-arCoeff()-X,1', 'tBodyAcc-arC
oeff()-X,2', 'tBodyAcc-arCoeff()-X,3', 'tBodyAcc-arCoeff()-X,4', 'tBody
Acc-arCoeff()-Y,1', 'tBodyAcc-arCoeff()-Y,2', 'tBodyAcc-arCoeff()-Y,3',
'tBodyAcc-arCoeff()-Y,4', 'tBodyAcc-arCoeff()-Z,1', 'tBodyAcc-arCoeff()
-Z,2', 'tBodyAcc-arCoeff()-Z,3', 'tBodyAcc-arCoeff()-Z,4', 'tBodyAcc-co
rrelation()-X,Y', 'tBodyAcc-correlation()-X,Z', 'tBodyAcc-correlation()
-Y,Z', 'tGravityAcc-mean()-X', 'tGravityAcc-mean()-Y', 'tGravityAcc-me
an()-Z', 'tGravityAcc-std()-X', 'tGravityAcc-std()-Y', 'tGravityAcc-st
d()-Z', 'tGravityAcc-mad()-X', 'tGravityAcc-mad()-Y', 'tGravityAcc-ma
d()-Z', 'tGravityAcc-max()-X', 'tGravityAcc-max()-Y', 'tGravityAcc-ma
x()-Z', 'tGravityAcc-min()-X', 'tGravityAcc-min()-Y', 'tGravityAcc-mi
n()-Z', 'tGravityAcc-sma()', 'tGravityAcc-energy()-X', 'tGravityAcc-ene
rgy()-Y', 'tGravityAcc-energy()-Z', 'tGravityAcc-iqr()-X', 'tGravityAcc
-iqr()-Y', 'tGravityAcc-iqr()-Z', 'tGravityAcc-entropy()-X', 'tGravityA
cc-entropy()-Y', 'tGravityAcc-entropy()-Z', 'tGravityAcc-arCoeff()-X,
1', 'tGravityAcc-arCoeff()-X,2', 'tGravityAcc-arCoeff()-X,3', 'tGravity
Acc-arCoeff()-X,4', 'tGravityAcc-arCoeff()-Y,1', 'tGravityAcc-arCoeff()
-Y,2', 'tGravityAcc-arCoeff()-Y,3', 'tGravityAcc-arCoeff()-Y,4', 'tGrav
ityAcc-arCoeff()-Z,1', 'tGravityAcc-arCoeff()-Z,2', 'tGravityAcc-arCoef
f()-Z,3', 'tGravityAcc-arCoeff()-Z,4', 'tGravityAcc-correlation()-X,Y',
'tGravityAcc-correlation()-X,Z', 'tGravityAcc-correlation()-Y,Z', 'tBod

```

yAccJerk-mean()-X', 'tBodyAccJerk-mean()-Y', 'tBodyAccJerk-mean()-Z', 'tBodyAccJerk-std()-X', 'tBodyAccJerk-std()-Y', 'tBodyAccJerk-std()-Z', 'tBodyAccJerk-mad()-X', 'tBodyAccJerk-mad()-Y', 'tBodyAccJerk-mad()-Z', 'tBodyAccJerk-max()-X', 'tBodyAccJerk-max()-Y', 'tBodyAccJerk-max()-Z', 'tBodyAccJerk-min()-X', 'tBodyAccJerk-min()-Y', 'tBodyAccJerk-min()-Z', 'tBodyAccJerk-sma()', 'tBodyAccJerk-energy()-X', 'tBodyAccJerk-energy()-Y', 'tBodyAccJerk-energy()-Z', 'tBodyAccJerk-iqr()-X', 'tBodyAccJerk-iqr()-Y', 'tBodyAccJerk-iqr()-Z', 'tBodyAccJerk-entropy()-X', 'tBodyAccJerk-entropy()-Y', 'tBodyAccJerk-entropy()-Z', 'tBodyAccJerk-arCoeff()-X,1', 'tBodyAccJerk-arCoeff()-X,2', 'tBodyAccJerk-arCoeff()-X,3', 'tBodyAccJerk-arCoeff()-X,4', 'tBodyAccJerk-arCoeff()-Y,1', 'tBodyAccJerk-arCoeff()-Y,2', 'tBodyAccJerk-arCoeff()-Y,3', 'tBodyAccJerk-arCoeff()-Y,4', 'tBodyAccJerk-arCoeff()-Z,1', 'tBodyAccJerk-arCoeff()-Z,2', 'tBodyAccJerk-arCoeff()-Z,3', 'tBodyAccJerk-arCoeff()-Z,4', 'tBodyAccJerk-corr elation()-X,Y', 'tBodyAccJerk-correlation()-X,Z', 'tBodyAccJerk-correla tion()-Y,Z', 'tBodyGyro-mean()-X', 'tBodyGyro-mean()-Y', 'tBodyGyro-me a()-Z', 'tBodyGyro-std()-X', 'tBodyGyro-std()-Y', 'tBodyGyro-std()-Z', 'tBodyGyro-mad()-X', 'tBodyGyro-mad()-Y', 'tBodyGyro-mad()-Z', 'tBodyGyro-max()-X', 'tBodyGyro-max()-Y', 'tBodyGyro-max()-Z', 'tBodyGyro-min()-X', 'tBodyGyro-min()-Y', 'tBodyGyro-min()-Z', 'tBodyGyro-sma()', 'tBodyGyro-energy()-X', 'tBodyGyro-energy()-Y', 'tBodyGyro-energy()-Z', 'tBo dyGyro-iqr()-X', 'tBodyGyro-iqr()-Y', 'tBodyGyro-iqr()-Z', 'tBodyGyro-e ntropy()-X', 'tBodyGyro-entropy()-Y', 'tBodyGyro-entropy()-Z', 'tBodyGyro- e-Entropy()-X,1', 'tBodyGyro-arCoeff()-X,2', 'tBodyGyro-arCoeff()-X,3', 'tBodyGyro-arCoeff()-X,4', 'tBodyGyro-arCoeff()-Y,1', 'tBodyGyro-arCoeff()-Y,2', 'tBodyGyro-arCoeff()-Y,3', 'tBodyGyro-arCoeff()-Y,4', 'tBodyGyro-arCoeff()-Z,1', 'tBodyGyro-arCoeff()-Z,2', 'tBodyGyro-arCoeff()-Z,3', 'tBodyGyro-arCoeff()-Z,4', 'tBodyGyro-correlation()-X,Y', 'tBodyGyro-correlation()-X,Z', 'tBodyGyroJerk-mean()-X', 'tBodyGyroJerk-mean()-Y', 'tBodyGyroJerk-mean()-Z', 'tBodyGyroJerk-std()-X', 'tBodyGyroJerk-std()-Y', 'tBodyGyroJerk-std()-Z', 'tBo dyGyroJerk-mad()-X', 'tBodyGyroJerk-mad()-Y', 'tBodyGyroJerk-mad()-Z', 'tBodyGyroJerk-max()-X', 'tBodyGyroJerk-max()-Y', 'tBodyGyroJerk-max()-Z', 'tBodyGyroJerk-min()-X', 'tBodyGyroJerk-min()-Y', 'tBodyGyroJerk-min()-Z', 'tBodyGyroJerk-sma()', 'tBodyGyroJerk-energy()-X', 'tBodyGyroJerk-energy()-Y', 'tBodyGyroJerk-energy()-Z', 'tBodyGyroJerk-iqr()-X', 'tBodyGyroJerk-iqr()-Y', 'tBodyGyroJerk-iqr()-Z', 'tBodyGyroJerk-entropy()-X', 'tBodyGyroJerk-entropy()-Y', 'tBodyGyroJerk-entropy()-Z', 'tBodyGyroJerk-arCoeff()-X,1', 'tBodyGyroJerk-arCoeff()-X,2', 'tBodyGyroJerk-arCoeff()-X,3', 'tBodyGyroJerk-arCoeff()-X,4', 'tBodyGyroJerk-arCoeff()-Y,1', 'tBodyGyroJerk-arCoeff()-Y,2', 'tBodyGyroJerk-arCoeff()-Y,3', 'tBodyGyroJerk-arCoeff()-Y,4', 'tBodyGyroJerk-arCoeff()-Z,1', 'tBodyGyroJerk-arCoeff()-Z,2', 'tBodyGyroJerk-arCoeff()-Z,3', 'tBodyGyroJerk-arCoeff()-Z,4', 'tBodyGyroJerk-correlation()-X,Y', 'tBodyGyroJerk-correlation()-X,Z', 'tBodyGyroJerk-correlation()-Y,Z', 'tBodyAccMag-mean()', 'tBodyAccMag-std()', 'tBodyAccMag-mad()', 'tBodyAccMag-max()', 'tBodyAccMag-min()', 'tBodyAccMag-sma()', 'tBodyAccMag-energy()', 'tBodyAccMag-iqr()', 'tBodyAccMag-entropy()', 'tBodyAccMag-arCoeff()1', 'tBodyAccMag-a rCoeff()2', 'tBodyAccMag-arCoeff()3', 'tBodyAccMag-arCoeff()4', 'tGravi tyAccMag-mean()', 'tGravityAccMag-std()', 'tGravityAccMag-mad()', 'tGr avityAccMag-max()', 'tGravityAccMag-min()', 'tGravityAccMag-sma()', 'tGr avityAccMag-energy()', 'tGravityAccMag-iqr()', 'tGravityAccMag-entrop

y()', 'tGravityAccMag-arCoeff()1', 'tGravityAccMag-arCoeff()2', 'tGravityAccMag-arCoeff()3', 'tGravityAccMag-arCoeff()4', 'tBodyAccJerkMag-mean()', 'tBodyAccJerkMag-std()', 'tBodyAccJerkMag-mad()', 'tBodyAccJerkMag-max()', 'tBodyAccJerkMag-min()', 'tBodyAccJerkMag-sma()', 'tBodyAccJerkMag-energy()', 'tBodyAccJerkMag-iqr()', 'tBodyAccJerkMag-entropy()', 'tBodyAccJerkMag-arCoeff()1', 'tBodyAccJerkMag-arCoeff()2', 'tBodyAccJerkMag-arCoeff()3', 'tBodyAccJerkMag-arCoeff()4', 'tBodyGyroMag-mean()', 'tBodyGyroMag-std()', 'tBodyGyroMag-mad()', 'tBodyGyroMag-max()', 'tBodyGyroMag-min()', 'tBodyGyroMag-sma()', 'tBodyGyroMag-energy()', 'tBodyGyroMag-iqr()', 'tBodyGyroMag-entropy()', 'tBodyGyroMag-arCoeff()1', 'tBodyGyroMag-arCoeff()2', 'tBodyGyroMag-arCoeff()3', 'tBodyGyroMag-arCoeff()4', 'tBodyGyroJerkMag-mean()', 'tBodyGyroJerkMag-std()', 'tBodyGyroJerkMag-mad()', 'tBodyGyroJerkMag-max()', 'tBodyGyroJerkMag-min()', 'tBodyGyroJerkMag-sma()', 'tBodyGyroJerkMag-energy()', 'tBodyGyroJerkMag-iqr()', 'tBodyGyroJerkMag-entropy()', 'tBodyGyroJerkMag-arCoeff()1', 'tBodyGyroJerkMag-arCoeff()2', 'tBodyGyroJerkMag-arCoeff()3', 'tBodyGyroJerkMag-arCoeff()4', 'fBodyAcc-mean()-X', 'fBodyAcc-mean()-Y', 'fBodyAcc-mean()-Z', 'fBodyAcc-mad()-X', 'fBodyAcc-mad()-Y', 'fBodyAcc-mad()-Z', 'fBodyAcc-max()-X', 'fBodyAcc-max()-Y', 'fBodyAcc-max()-Z', 'fBodyAcc-min()-X', 'fBodyAcc-min()-Y', 'fBodyAcc-min()-Z', 'fBodyAcc-sma()', 'fBodyAcc-energy()-X', 'fBodyAcc-energy()-Y', 'fBodyAcc-energy()-Z', 'fBodyAcc-iqr()-X', 'fBodyAcc-iqr()-Y', 'fBodyAcc-iqr()-Z', 'fBodyAcc-entropy()-X', 'fBodyAcc-entropy()-Y', 'fBodyAcc-entropy()-Z', 'fBodyAcc-maxInds-X', 'fBodyAcc-maxInds-Y', 'fBodyAcc-maxInds-Z', 'fBodyAcc-meanFreq()-X', 'fBodyAcc-meanFreq()-Y', 'fBodyAcc-meanFreq()-Z', 'fBodyAcc-skewness()-X', 'fBodyAcc-kurtosis()-X', 'fBodyAcc-skewness()-Y', 'fBodyAcc-kurtosis()-Y', 'fBodyAcc-skewness()-Z', 'fBodyAcc-kurtosis()-Z', 'fBodyAcc-bandsEnergy()-1,8', 'fBodyAcc-bandsEnergy()-9,16', 'fBodyAcc-bandsEnergy()-17,24', 'fBodyAcc-bandsEnergy()-25,32', 'fBodyAcc-bandsEnergy()-33,40', 'fBodyAcc-bandsEnergy()-41,48', 'fBodyAcc-bandsEnergy()-49,56', 'fBodyAcc-bandsEnergy()-57,64', 'fBodyAcc-bandsEnergy()-1,16', 'fBodyAcc-bandsEnergy()-17,32', 'fBodyAcc-bandsEnergy()-33,48', 'fBodyAcc-bandsEnergy()-49,64', 'fBodyAcc-bandsEnergy()-1,24', 'fBodyAcc-bandsEnergy()-25,48', 'fBodyAcc-bandsEnergy()-1,8.1', 'fBodyAcc-bandsEnergy()-9,16.1', 'fBodyAcc-bandsEnergy()-17,24.1', 'fBodyAcc-bandsEnergy()-25,32.1', 'fBodyAcc-bandsEnergy()-33,40.1', 'fBodyAcc-bandsEnergy()-41,48.1', 'fBodyAcc-bandsEnergy()-49,56.1', 'fBodyAcc-bandsEnergy()-57,64.1', 'fBodyAcc-bandsEnergy()-1,16.1', 'fBodyAcc-bandsEnergy()-17,32.1', 'fBodyAcc-bandsEnergy()-33,48.1', 'fBodyAcc-bandsEnergy()-49,64.1', 'fBodyAcc-bandsEnergy()-1,24.1', 'fBodyAcc-bandsEnergy()-25,48.1', 'fBodyAcc-bandsEnergy()-1,8.2', 'fBodyAcc-bandsEnergy()-9,16.2', 'fBodyAcc-bandsEnergy()-17,24.2', 'fBodyAcc-bandsEnergy()-25,32.2', 'fBodyAcc-bandsEnergy()-33,40.2', 'fBodyAcc-bandsEnergy()-41,48.2', 'fBodyAcc-bandsEnergy()-49,56.2', 'fBodyAcc-bandsEnergy()-57,64.2', 'fBodyAcc-bandsEnergy()-1,16.2', 'fBodyAcc-bandsEnergy()-17,32.2', 'fBodyAcc-bandsEnergy()-33,48.2', 'fBodyAcc-bandsEnergy()-49,64.2', 'fBodyAcc-bandsEnergy()-1,24.2', 'fBodyAcc-bandsEnergy()-25,48.2', 'fBodyAccJerk-mean()-X', 'fBodyAccJerk-mean()-Y', 'fBodyAccJerk-mean()-Z', 'fBodyAccJerk-std()-X', 'fBodyAccJerk-std()-Y', 'fBodyAccJerk-std()-Z', 'fBodyAccJerk-mad()-X', 'fBodyAccJerk-mad()-Y', 'fBodyAccJerk-mad()-Z', 'fBodyAccJerk-max()-X', 'fBodyAccJerk-max()-Y', 'fBodyAccJerk-max()-Z', 'fBodyAccJerk-min()-X', 'fBodyAccJerk-min()-Y', 'fBodyAccJerk-min()-Z'

'fBodyAccJerk-min()-Z', 'fBodyAccJerk-sma()', 'fBodyAccJerk-energy()-X', 'fBodyAccJerk-energy()-Y', 'fBodyAccJerk-energy()-Z', 'fBodyAccJerk-iqr()-X', 'fBodyAccJerk-iqr()-Y', 'fBodyAccJerk-iqr()-Z', 'fBodyAccJerk-entropy()-X', 'fBodyAccJerk-entropy()-Y', 'fBodyAccJerk-entropy()-Z', 'fBodyAccJerk-maxInds-X', 'fBodyAccJerk-maxInds-Y', 'fBodyAccJerk-maxInds-Z', 'fBodyAccJerk-meanFreq()-X', 'fBodyAccJerk-meanFreq()-Y', 'fBodyAccJerk-meanFreq()-Z', 'fBodyAccJerk-skewness()-X', 'fBodyAccJerk-kurtosis()-X', 'fBodyAccJerk-skewness()-Y', 'fBodyAccJerk-kurtosis()-Y', 'fBodyAccJerk-bandsEnergy()-1,8', 'fBodyAccJerk-bandsEnergy()-9,16', 'fBodyAccJerk-bandsEnergy()-17,24', 'fBodyAccJerk-bandsEnergy()-25,32', 'fBodyAccJerk-bandsEnergy()-33,40', 'fBodyAccJerk-bandsEnergy()-41,48', 'fBodyAccJerk-bandsEnergy()-49,56', 'fBodyAccJerk-bandsEnergy()-57,64', 'fBodyAccJerk-bandsEnergy()-1,16', 'fBodyAccJerk-bandsEnergy()-17,32', 'fBodyAccJerk-bandsEnergy()-33,48', 'fBodyAccJerk-bandsEnergy()-49,64', 'fBodyAccJerk-bandsEnergy()-1,24', 'fBodyAccJerk-bandsEnergy()-25,48', 'fBodyAccJerk-bandsEnergy()-1,8.1', 'fBodyAccJerk-bandsEnergy()-9,16.1', 'fBodyAccJerk-bandsEnergy()-17,24.1', 'fBodyAccJerk-bandsEnergy()-25,32.1', 'fBodyAccJerk-bandsEnergy()-33,40.1', 'fBodyAccJerk-bandsEnergy()-41,48.1', 'fBodyAccJerk-bandsEnergy()-49,56.1', 'fBodyAccJerk-bandsEnergy()-57,64.1', 'fBodyAccJerk-bandsEnergy()-1,16.1', 'fBodyAccJerk-bandsEnergy()-17,32.1', 'fBodyAccJerk-bandsEnergy()-33,48.1', 'fBodyAccJerk-bandsEnergy()-49,64.1', 'fBodyAccJerk-bandsEnergy()-1,24.1', 'fBodyAccJerk-bandsEnergy()-25,48.1', 'fBodyAccJerk-bandsEnergy()-9,16.2', 'fBodyAccJerk-bandsEnergy()-17,24.2', 'fBodyAccJerk-bandsEnergy()-25,32.2', 'fBodyAccJerk-bandsEnergy()-33,40.2', 'fBodyAccJerk-bandsEnergy()-41,48.2', 'fBodyAccJerk-bandsEnergy()-49,56.2', 'fBodyAccJerk-bandsEnergy()-57,64.2', 'fBodyAccJerk-bandsEnergy()-1,16.2', 'fBodyAccJerk-bandsEnergy()-17,32.2', 'fBodyAccJerk-bandsEnergy()-33,48.2', 'fBodyAccJerk-bandsEnergy()-49,64.2', 'fBodyAccJerk-bandsEnergy()-1,24.2', 'fBodyAccJerk-bandsEnergy()-25,48.2', 'fBodyGyro-mean()-X', 'fBodyGyro-mean()-Y', 'fBodyGyro-mean()-Z', 'fBodyGyro-std()-X', 'fBodyGyro-std()-Y', 'fBodyGyro-std()-Z', 'fBodyGyro-mad()-X', 'fBodyGyro-mad()-Y', 'fBodyGyro-mad()-Z', 'fBodyGyro-max()-X', 'fBodyGyro-max()-Y', 'fBodyGyro-max()-Z', 'fBodyGyro-min()-X', 'fBodyGyro-min()-Y', 'fBodyGyro-min()-Z', 'fBodyGyro-sma()', 'fBodyGyro-energy()-X', 'fBodyGyro-energy()-Y', 'fBodyGyro-energy()-Z', 'fBodyGyro-iqr()-X', 'fBodyGyro-iqr()-Y', 'fBodyGyro-iqr()-Z', 'fBodyGyro-entropy()-X', 'fBodyGyro-entropy()-Y', 'fBodyGyro-entropy()-Z', 'fBodyGyro-maxInds-X', 'fBodyGyro-maxInds-Y', 'fBodyGyro-maxInds-Z', 'fBodyGyro-meanFreq()-X', 'fBodyGyro-meanFreq()-Y', 'fBodyGyro-meanFreq()-Z', 'fBodyGyro-skewness()-X', 'fBodyGyro-kurtosis()-X', 'fBodyGyro-skewness()-Y', 'fBodyGyro-kurtosis()-Y', 'fBodyGyro-bandsEnergy()-1,8', 'fBodyGyro-bandsEnergy()-9,16', 'fBodyGyro-bandsEnergy()-17,24', 'fBodyGyro-bandsEnergy()-25,32', 'fBodyGyro-bandsEnergy()-41,48', 'fBodyGyro-bandsEnergy()-49,56', 'fBodyGyro-bandsEnergy()-57,64', 'fBodyGyro-bandsEnergy()-1,16', 'fBodyGyro-bandsEnergy()-17,32', 'fBodyGyro-bandsEnergy()-33,48', 'fBodyGyro-bandsEnergy()-49,64', 'fBodyGyro-bandsEnergy()-1,24', 'fBodyGyro-bandsEnergy()-25,48', 'fBodyGyro-bandsEnergy()-1,8.1', 'fBodyGyro-bandsEnergy()-9,16.1', 'fBodyGyro-bandsEnergy()-17,24.1', 'fBodyGyro-bandsEnergy()-25,32.1', 'fBodyGyro-bandsEnergy()-33,40.1', 'fBodyGyro-bandsEnergy()-41,48.1', 'fBodyGyro-bandsEnergy()-49,56.1', 'fBodyGyro-bandsEnergy()-57,64.1'

```
1,48.1', 'fBodyGyro-bandsEnergy()-49,56.1', 'fBodyGyro-bandsEnergy()-5
7,64.1', 'fBodyGyro-bandsEnergy()-1,16.1', 'fBodyGyro-bandsEnergy()-17,
32.1', 'fBodyGyro-bandsEnergy()-33,48.1', 'fBodyGyro-bandsEnergy()-49,6
4.1', 'fBodyGyro-bandsEnergy()-1,24.1', 'fBodyGyro-bandsEnergy()-25,48.
1', 'fBodyGyro-bandsEnergy()-1,8.2', 'fBodyGyro-bandsEnergy()-9,16.2',
'fBodyGyro-bandsEnergy()-17,24.2', 'fBodyGyro-bandsEnergy()-25,32.2', 'f
BodyGyro-bandsEnergy()-33,40.2', 'fBodyGyro-bandsEnergy()-41,48.2', 'f
BodyGyro-bandsEnergy()-49,56.2', 'fBodyGyro-bandsEnergy()-57,64.2', 'fB
odyGyro-bandsEnergy()-1,16.2', 'fBodyGyro-bandsEnergy()-17,32.2', 'fBod
yGyro-bandsEnergy()-33,48.2', 'fBodyGyro-bandsEnergy()-49,64.2', 'fBody
Gyro-bandsEnergy()-1,24.2', 'fBodyGyro-bandsEnergy()-25,48.2', 'fBodyAc
cMag-mean()', 'fBodyAccMag-std()', 'fBodyAccMag-mad()', 'fBodyAccMag-ma
x()', 'fBodyAccMag-min()', 'fBodyAccMag-sma()', 'fBodyAccMag-energy()', 'f
BodyAccMag-iqr()', 'fBodyAccMag-entropy()', 'fBodyAccMag-maxInds', 'f
BodyAccMag-meanFreq()', 'fBodyAccMag-skewness()', 'fBodyAccMag-kurtosi
s()', 'fBodyBodyAccJerkMag-mean()', 'fBodyBodyAccJerkMag-std()', 'fBody
BodyAccJerkMag-mad()', 'fBodyBodyAccJerkMag-max()', 'fBodyBodyAccJerkMa
g-min()', 'fBodyBodyAccJerkMag-sma()', 'fBodyBodyAccJerkMag-energy()', 'f
BodyBodyAccJerkMag-iqr()', 'fBodyBodyAccJerkMag-entropy()', 'fBodyBod
yAccJerkMag-maxInds', 'fBodyBodyAccJerkMag-meanFreq()', 'fBodyBodyAccJe
rkMag-skewness()', 'fBodyBodyAccJerkMag-kurtosis()', 'fBodyBodyGyroMag-
mean()', 'fBodyBodyGyroMag-std()', 'fBodyBodyGyroMag-mad()', 'fBodyBody
GyroMag-max()', 'fBodyBodyGyroMag-min()', 'fBodyBodyGyroMag-sma()', 'fB
odyBodyGyroMag-energy()', 'fBodyBodyGyroMag-iqr()', 'fBodyBodyGyroMag-e
ntropy()', 'fBodyBodyGyroMag-maxInds', 'fBodyBodyGyroMag-meanFreq()', 'f
BodyBodyGyroMag-skewness()', 'fBodyBodyGyroMag-kurtosis()', 'fBodyBody
GyroJerkMag-mean()', 'fBodyBodyGyroJerkMag-std()', 'fBodyBodyGyroJerkMa
g-mad()', 'fBodyBodyGyroJerkMag-max()', 'fBodyBodyGyroJerkMag-min()', 'f
BodyBodyGyroJerkMag-sma()', 'fBodyBodyGyroJerkMag-energy()', 'fBodyBod
yGyroJerkMag-iqr()', 'fBodyBodyGyroJerkMag-entropy()', 'fBodyBodyGyroJe
rkMag-maxInds', 'fBodyBodyGyroJerkMag-meanFreq()', 'fBodyBodyGyroJerkMa
g-skewness()', 'fBodyBodyGyroJerkMag-kurtosis()', 'angle(tBodyAccMean,g
ravity)', 'angle(tBodyAccJerkMean),gravityMean)', 'angle(tBodyGyroMean,
gravityMean)', 'angle(tBodyGyroJerkMean,gravityMean)', 'angle(X,gravity
Mean)', 'angle(Y,gravityMean)', 'angle(Z,gravityMean)', 'subject', 'Act
ivity']
```

```
In [3]: # ----- 2. Data Exploration -----
print("\n== Basic info ==")
print(df.info())

total_missing = df.isna().sum().sum()
print("\nTotal number of missing values in the entire dataset: ", tota
missing_per_column = df.isna().sum()
print("\nMissing values per column (only columns with missing values):")
print(missing_per_column[missing_per_column > 0])

print("\n== Numerical Summary Statistics for the first 10 features ==")
print(df.describe().T.head(10))

print("\n== Activity Label Distribution ==")
activity_counts = df["Activity"].value_counts()
```

```

print(activity_counts)

plt.figure(figsize=(10, 10))
sns.countplot(x="Activity", data=df, order=activity_counts.index)
plt.title("Activity Label Distribution")
plt.tight_layout()
plt.show()

print("\n== Number of samples per subject ==")
subject_count = df["subject"].value_counts().sort_index()
print(subject_count.head(10))

plt.figure(figsize=(10, 10))
subject_count.plot(kind="bar")
plt.title("Number of Samples per Subject")
plt.xlabel("Subject ID")
plt.ylabel("Count")
plt.tight_layout()
plt.show()

# 2.5 Look for Outliers and key features
features_to_inspect = ["tBodyAcc-mean()-X", "tBodyAcc-mean()-Y", "tBodyAcc-mean()-Z", "tBodyAcc-std()-X", "tBodyAcc-std()-Y", "tBodyAcc-std()-Z", "tBodyGyro-mean()-X", "tBodyGyro-mean()-Y", "tBodyGyro-mean()-Z", "tBodyGyro-std()-X", "tBodyGyro-std()-Y", "tBodyGyro-std()-Z", "tBodyAccJerk-mean()-X", "tBodyAccJerk-mean()-Y", "tBodyAccJerk-mean()-Z", "tBodyAccJerk-std()-X", "tBodyAccJerk-std()-Y", "tBodyAccJerk-std()-Z", "tBodyGyroJerk-mean()-X", "tBodyGyroJerk-mean()-Y", "tBodyGyroJerk-mean()-Z", "tBodyGyroJerk-std()-X", "tBodyGyroJerk-std()-Y", "tBodyGyroJerk-std()-Z", "tBodyGyroZ-mean()", "tBodyGyroZ-std()", "tBodyAccJerkZ-mean()", "tBodyAccJerkZ-std()", "tBodyGyroJerkZ-mean()", "tBodyGyroJerkZ-std()", "tBodyGyroZ-std()", "tBodyGyroJerkZ-std()"]
print("\n== Example Features to inspect ==")
print(features_to_inspect)

fig, axes = plt.subplots(2, 2, figsize=(8, 8))
for ax, col in zip(axes.ravel(), features_to_inspect):
    ax.hist(df[col], bins=30)
    ax.set_title(col)
fig.suptitle("Histograms of selected features", y=1.02)
fig.tight_layout()
plt.show()

fig, ax = plt.subplots(figsize=(8, 6))
sns.boxplot(data=df[features_to_inspect], ax=ax)
ax.set_title("Boxplot of selected features (outlier inspection)")
ax.set_ylabel("Standardized measurement units")
ax.tick_params(axis='x', rotation=20)
fig.tight_layout()
plt.show()

corr_subset = df[features_to_inspect].corr()
print("\n== Correlation Matrix for the selected features==")
print(corr_subset)

fig, ax = plt.subplots(figsize=(8, 8))
sns.heatmap(corr_subset, annot=True, fmt=".2f", cmap="coolwarm", square=True)
ax.set_title("Correlation HeatMap for selected features")
fig.tight_layout()
plt.show()

== Basic info ==
<class 'pandas.core.frame.DataFrame'>

```

```
RangeIndex: 7352 entries, 0 to 7351
Columns: 563 entries, tBodyAcc-mean()-X to Activity
dtypes: float64(561), int64(1), object(1)
memory usage: 31.6+ MB
None
```

```
Total number of missing values in the entire dataset: 0
```

```
Missing values per column (only columns with missing values):
Series([], dtype: int64)
```


```
== Numerical Summary Statistics for the first 10 features ==
      count      mean      std      min      25%
50% \
tBodyAcc-mean()-X 7352.0  0.274488  0.070261 -1.000000  0.262975  0.27
7193
tBodyAcc-mean()-Y 7352.0 -0.017695  0.040811 -1.000000 -0.024863 -0.01
7219
tBodyAcc-mean()-Z 7352.0 -0.109141  0.056635 -1.000000 -0.120993 -0.10
8676
tBodyAcc-std()-X 7352.0 -0.605438  0.448734 -1.000000 -0.992754 -0.94
6196
tBodyAcc-std()-Y 7352.0 -0.510938  0.502645 -0.999873 -0.978129 -0.85
1897
tBodyAcc-std()-Z 7352.0 -0.604754  0.418687 -1.000000 -0.980233 -0.85
9365
tBodyAcc-mad()-X 7352.0 -0.630512  0.424073 -1.000000 -0.993591 -0.95
0709
tBodyAcc-mad()-Y 7352.0 -0.526907  0.485942 -1.000000 -0.978162 -0.85
7328
tBodyAcc-mad()-Z 7352.0 -0.606150  0.414122 -1.000000 -0.980251 -0.85
7143
tBodyAcc-max()-X 7352.0 -0.468604  0.544547 -1.000000 -0.936219 -0.88
1637
```

```
      75%      max
tBodyAcc-mean()-X 0.288461  1.000000
tBodyAcc-mean()-Y -0.010783  1.000000
tBodyAcc-mean()-Z -0.097794  1.000000
tBodyAcc-std()-X -0.242813  1.000000
tBodyAcc-std()-Y -0.034231  0.916238
tBodyAcc-std()-Z -0.262415  1.000000
tBodyAcc-mad()-X -0.292680  1.000000
tBodyAcc-mad()-Y -0.066701  0.967664
tBodyAcc-mad()-Z -0.265671  1.000000
tBodyAcc-max()-X -0.017129  1.000000
```

```
== Activity Label Distribution ==
```

Activity	
LAYING	1407
STANDING	1374
SITTING	1286


```
WALKING           1226
WALKING_UPSTAIRS  1073
WALKING_DOWNSTAIRS 986
Name: count, dtype: int64
```



```
== Number of samples per subject ==
subject
```

```
1      347
3      341
5      302
6      325
7      308
8      281
11     316
14     323
15     328
16     366
```


```
Name: count, dtype: int64
```




== Example Features to inspect ==

```
['tBodyAcc-mean()-X', 'tBodyAcc-mean()-Y', 'tBodyAcc-std()-X', 'tBodyAcc-std()-Y']
```


Histograms of selected features





== Correlation Matrix for the selected features==

	tBodyAcc-mean()-X	tBodyAcc-mean()-Y	tBodyAcc-std()
-X \			
tBodyAcc-mean()-X	1.000000	0.148061	0.0006
19			
tBodyAcc-mean()-Y	0.148061	1.000000	-0.0451
60			
tBodyAcc-std()-X	0.000619	-0.045160	1.0000
00			
tBodyAcc-std()-Y	-0.021903	-0.044920	0.9274
61			
		tBodyAcc-std()-Y	
tBodyAcc-mean()-X	-0.021903		
tBodyAcc-mean()-Y	-0.044920		
tBodyAcc-std()-X	0.927461		
tBodyAcc-std()-Y	1.000000		



```
In [4]: # ----- 3. Prepare Features / Split -----
# Drop ID-like column to avoid leakage
X = df.drop(columns=['Activity', 'subject'])
y = df['Activity']
num_features = X.columns.tolist()

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
print("\n== Train/Test Sizes==")
print(f"X_train: {X_train.shape}, y_train: {y_train.shape}")

# Preprocessing
scaler = ColumnTransformer(transformers=[("num", StandardScaler(), num

def scaled_pipeline(estimator):
    return Pipeline(steps=[("scaler", scaler), ("estimator", estimator
```

```
def passthrough_pipeline(estimator):
    return Pipeline(steps=[("clf", estimator)])  
  
== Train/Test Sizes==  
X_train: (5881, 561), y_train: (5881,)
```

```
In [5]: # ----- 3.4 Models -----
models = {
    "Perceptron": scaled_pipeline(Perceptron(random_state=42)),
    "RidgeClassifier (Adaline-like)": scaled_pipeline(RidgeClassifier()),
    "Logistic Regression": scaled_pipeline(LogisticRegression(max_iter=1000)),
    "SVM (RBF)": scaled_pipeline(SVC(kernel="rbf", C=10, gamma="scale")),
    "KNN (k=15)": scaled_pipeline(KNeighborsClassifier(n_neighbors=15)),
    "DecisionTree": passthrough_pipeline(DecisionTreeClassifier(max_depth=5)),
    "RandomForest": passthrough_pipeline(RandomForestClassifier(n_estimators=100))
}
```

```
In [6]: # 1. Importing all the important libraries
import warnings
warnings.filterwarnings("ignore")  
  
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns  
  
from sklearn.model_selection import train_test_split, StratifiedKFold
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer  
  
from sklearn.linear_model import Perceptron, RidgeClassifier, LogisticRegression
from sklearn.svm import SVC
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, f1_score, classification_report  
  
from sklearn.base import clone
from scipy.stats import loguniform, randint as sp_randint
from joblib import dump
```

```
In [7]: # ----- 3.5 Train & Evaluate -----
results = []
reports = {}  
  
for name, pipe in models.items():
    pipe.fit(X_train, y_train)
    y_pred = pipe.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    f1m = f1_score(y_test, y_pred, average="macro")
    results.append({"model": name, "accuracy": acc, "f1_macro": f1m})
```

```
reports[name] = classification_report(y_test, y_pred, digits=3)
print(f"\n== {name} ==")
print(f"Accuracy: {acc:.4f} | f1_macro: {f1m:.4f}")

results_df = pd.DataFrame(results).sort_values(by=["accuracy", "f1_macro"], ascending=False)
print("\n== Model Leaderboard (Test Set) ==")
print(results_df.to_string(index=False))

best_name = results_df.iloc[0]["model"]
print(f"\n== Best Model: {best_name} ==")
print(reports[best_name])

best_pipe = models[best_name]
y_pred_best = best_pipe.predict(X_test)
fig, ax = plt.subplots(figsize=(8, 6))
ConfusionMatrixDisplay.from_predictions(
    y_test, y_pred_best, xticks_rotation=45, cmap="Blues", colorbar=False
)
ax.set_title(f"Confusion Matrix for {best_name}")
fig.tight_layout()
plt.show()
```

```

== Perceptron ==
Accuracy: 0.9810 | f1_macro: 0.9818

== RidgeClassifier (Adaline-like) ==
Accuracy: 0.9810 | f1_macro: 0.9822

== Logistic Regression ==
Accuracy: 0.9857 | f1_macro: 0.9868

== SVM (RBF) ==
Accuracy: 0.9898 | f1_macro: 0.9906

== KNN (k=15) ==
Accuracy: 0.9565 | f1_macro: 0.9578

== DecisionTree ==
Accuracy: 0.9388 | f1_macro: 0.9380


== RandomForest ==
Accuracy: 0.9878 | f1_macro: 0.9879

== Model Leaderboard (Test Set) ==
      model  accuracy  f1_macro
      SVM (RBF)  0.989803  0.990587
      RandomForest  0.987763  0.987936
      Logistic Regression  0.985724  0.986825
      RidgeClassifier (Adaline-like)  0.980965  0.982162
      Perceptron  0.980965  0.981771
      KNN (k=15)  0.956492  0.957820
      DecisionTree  0.938817  0.937979

== Best Model: SVM (RBF) ==
      precision    recall  f1-score   support
      LAYING      1.000     1.000     1.000      282
      SITTING     0.976     0.965     0.971      257
      STANDING     0.968     0.978     0.973      275
      WALKING      1.000     1.000     1.000      245
      WALKING_DOWNSTAIRS  1.000     1.000     1.000      197
      WALKING_UPSTAIRS   1.000     1.000     1.000      215

      accuracy          0.990      1471
      macro avg       0.991      0.991      1471
      weighted avg    0.990      0.990      1471

```



```
In [8]: # ----- 4. Hyperparameter Search -----
# ----- Runtime knobs -----
FAST_SEARCH = True # True = faster hyperparam search; set False for
cv_splits = 3 if FAST_SEARCH else 5
cv = StratifiedKFold(n_splits=cv_splits, shuffle=True, random_state=42

def eval_and_log(name, estimator, Xtr=X_train, ytr=y_train, xte=X_test):
    estimator = estimator.fit(Xtr, ytr)
    yhat = estimator.predict(xte)
    acc = accuracy_score(yte, yhat)
    f1m = f1_score(yte, yhat, average="macro")
    print(f"\n== {name} ==")
    print(f"Accuracy: {acc:.4f} | f1_macro: {f1m:.4f}")
    return {"model": name, "estimator": estimator, "accuracy": acc, "f
```

Parameter spaces (target the inner estimator in each pipeline)

```
search_specs = {
    "SVM (RBF)": {
        "pipe": models["SVM (RBF)"],
```

```

    "params": {
        "estimator__C": loguniform(1e-1, 1e2),    # 0.1-100
        "estimator__gamma": ["scale", "auto"]
    },
    "n_iter": 15 if FAST_SEARCH else 25
},
"RandomForest": {
    "pipe": models["RandomForest"],
    "params": {
        "clf__n_estimators": sp_randint(150, 401),  # 150-400 tree
        "clf__max_depth": [None, 10, 20, 30],
        "clf__min_samples_split": sp_randint(2, 11),
        "clf__min_samples_leaf": sp_randint(1, 5),
        "clf__max_features": ["sqrt", "log2", None]
    },
    "n_iter": 15 if FAST_SEARCH else 30
},
"Logistic Regression": {
    "pipe": models["Logistic Regression"],
    "params": {
        "estimator__C": loguniform(1e-2, 1e2),
        "estimator__penalty": ["l2"],
        "estimator__solver": ["lbfgs", "saga"]
    },
    "n_iter": 15 if FAST_SEARCH else 25
},
"KNN (k=15)": {
    "pipe": models["KNN (k=15)"],
    "params": {
        "estimator__n_neighbors": sp_randint(5, 31),
        "estimator__weights": ["uniform", "distance"],
        "estimator__p": [1, 2]
    },
    "n_iter": 15 if FAST_SEARCH else 25
},
}

print("\n == Hyperparameter search (RandomizedSearchCV) ==")
tuned_models = {}
for name, spec in search_specs.items():
    search = RandomizedSearchCV(
        estimator=spec["pipe"],
        param_distributions=spec["params"],
        n_iter=spec["n_iter"],
        cv=cv,
        scoring="f1_macro",
        random_state=42,
        n_jobs=-1,
        verbose=1,
        refit=True,
        error_score="raise"
    )

```

```
search.fit(X_train, y_train)
print(f"\n == Best params for {name}: {search.best_params_}")
print(f"CV best f1_macro: {search.best_score_:.4f}")
tuned_models[f"{name} (tuned)"] = search.best_estimator_

# Evaluate tuned models on the test set and extend the leaderboard
more_results = []
for nm, est in tuned_models.items():
    more_results.append(eval_and_log(nm, est))

if more_results:
    more_df = pd.DataFrame([{k: v for k, v in r.items() if k != "estim
    results_df = (
        pd.concat([results_df, more_df[["model", "accuracy", "f1_macro"
            .sort_values(by=["accuracy", "f1_macro"], ascending=False)
    )
    print("\n == Updated Leaderboard (incl. tuned) ==")
    print(results_df.to_string(index=False))
```

```

== Hyperparameter search (RandomizedSearchCV) ==
Fitting 3 folds for each of 15 candidates, totalling 45 fits

== Best params for SVM (RBF): {'estimator__C': np.float64(31.428808908
401084), 'estimator__gamma': 'auto'}
CV best f1_macro: 0.9830
Fitting 3 folds for each of 15 candidates, totalling 45 fits

== Best params for RandomForest: {'clf__max_depth': 30, 'clf__max_features': 'log2', 'clf__min_samples_leaf': 2, 'clf__min_samples_split': 3, 'clf__n_estimators': 341}
CV best f1_macro: 0.9757
Fitting 3 folds for each of 15 candidates, totalling 45 fits

== Best params for Logistic Regression: {'estimator__C': np.float64(0.
5342937261279778), 'estimator__penalty': 'l2', 'estimator__solver': 'lb
fgs'}
CV best f1_macro: 0.9812
Fitting 3 folds for each of 15 candidates, totalling 45 fits

== Best params for KNN (k=15): {'estimator__n_neighbors': 16, 'estimat
or_p': 1, 'estimator__weights': 'uniform'}
CV best f1_macro: 0.9626

== SVM (RBF) (tuned) ==
Accuracy: 0.9898 | f1_macro: 0.9906

== RandomForest (tuned) ==
Accuracy: 0.9850 | f1_macro: 0.9853

== Logistic Regression (tuned) ==
Accuracy: 0.9857 | f1_macro: 0.9868

== KNN (k=15) (tuned) ==
Accuracy: 0.9708 | f1_macro: 0.9721

== Updated Leaderboard (incl. tuned) ==
      model  accuracy  f1_macro
      SVM (RBF) (tuned)  0.989803  0.990592
      SVM (RBF)          0.989803  0.990587
      RandomForest        0.987763  0.987936
      Logistic Regression 0.985724  0.986825
      Logistic Regression (tuned)  0.985724  0.986825
      RandomForest (tuned)  0.985044  0.985288
      RidgeClassifier (Adaline-like) 0.980965  0.982162
      Perceptron          0.980965  0.981771
      KNN (k=15) (tuned)   0.970768  0.972094
      KNN (k=15)          0.956492  0.957820
      DecisionTree         0.938817  0.937979

```

In [9]: # ----- 5. True Ensembles: Voting + Stacking -----

```

def ensure_svc_probability(pipe):
    p = clone(pipe)

```

```

        if "estimator" in p.named_steps and isinstance(p.named_steps["estimator"], LogisticRegression):
            p.set_params(estimator__probability=True)
    return p

svm_base = tuned_models.get("SVM (RBF) (tuned)", models["SVM (RBF)"])
rf_base = tuned_models.get("RandomForest (tuned)", models["RandomForest"])
lr_base = tuned_models.get("Logistic Regression (tuned)", models["Logistic Regression"])
knn_base = tuned_models.get("KNN (k=15) (tuned)", models["KNN (k=15)"])

svm_soft = ensure_svc_probability(svm_base)

voting_hard = VotingClassifier(
    estimators=[
        ("perc", models["Perceptron"]),
        ("ridge", models["RidgeClassifier (Adaline-like)"]),
        ("lr", lr_base),
        ("svm", svm_base),
        ("dt", models["DecisionTree"]),
        ("rf", rf_base),
        ("knn", knn_base),
    ],
    voting="hard",
    n_jobs=-1
)

voting_soft = VotingClassifier(
    estimators=[
        ("lr", lr_base),
        ("svm", svm_soft),
        ("rf", rf_base),
        ("knn", knn_base),
    ],
    voting="soft",
    n_jobs=-1
)

stacking = StackingClassifier(
    estimators=[
        ("lr", lr_base),
        ("svm", svm_soft),
        ("rf", rf_base),
        ("knn", knn_base),
    ],
    final_estimator=LogisticRegression(max_iter=5000, random_state=42),
    stack_method="predict_proba",
    cv=cv,
    n_jobs=-1,
    passthrough=False
)

ensemble_models = {

```

```

        "Voting (hard)": voting_hard,
        "Voting (soft)": voting_soft,
        "Stacking (LR meta)": stacking,
    }

print("\n== Ensembles ==")
ensemble_results = []
for nm, est in ensemble_models.items():
    ensemble_results.append(eval_and_log(nm, est))

== Ensembles ==

== Voting (hard) ==
Accuracy: 0.9905 | f1_macro: 0.9912

== Voting (soft) ==
Accuracy: 0.9918 | f1_macro: 0.9925

== Stacking (LR meta) ==
Accuracy: 0.9946 | f1_macro: 0.9950

```

```

In [10]: # ----- 6. Consolidate and choose global best -----
all_results = []
for r in results:
    all_results.append({"model": r["model"], "accuracy": r["accuracy"]}
for r in more_results:
    all_results.append({"model": r["model"], "accuracy": r["accuracy"]}
for r in ensemble_results:
    all_results.append({"model": r["model"], "accuracy": r["accuracy"]}

all_df = pd.DataFrame(all_results).sort_values(by=["accuracy", "f1_mac
print("\n == GRAND LEADERBOARD ==")
print(all_df.to_string(index=False))

best_overall_name = all_df.iloc[0]["model"]
print(f"\n== Best Overall: {best_overall_name} ==")

name_to_est = {**{k: v for k, v in models.items()}, **tuned_models, **}
best_overall_est = name_to_est[best_overall_name]
best_overall_est.fit(X_train, y_train)

fig, ax = plt.subplots(figsize=(8, 6))
ConfusionMatrixDisplay.from_predictions(
    y_test, best_overall_est.predict(X_test), xticks_rotation=45, cmap
)
ax.set_title(f"Confusion Matrix for {best_overall_name}")
fig.tight_layout()
plt.show()


```

== GRAND LEADERBOARD ==

	model	accuracy	f1_macro
Stacking (LR meta)	0.994562	0.994980	
Voting (soft)	0.991842	0.992473	
Voting (hard)	0.990483	0.991216	
SVM (RBF) (tuned)	0.989803	0.990592	
SVM (RBF)	0.989803	0.990587	
RandomForest	0.987763	0.987936	
Logistic Regression	0.985724	0.986825	
Logistic Regression (tuned)	0.985724	0.986825	
RandomForest (tuned)	0.985044	0.985288	
RidgeClassifier (Adaline-like)	0.980965	0.982162	
Perceptron	0.980965	0.981771	
KNN (k=15) (tuned)	0.970768	0.972094	
KNN (k=15)	0.956492	0.957820	
DecisionTree	0.938817	0.937979	

== Best Overall: Stacking (LR meta) ==

Confusion Matrix for Stacking (LR meta)



```
In [11]: # ----- 7. Persist the final pipeline & inference
MODEL_PATH = "best_har_model.joblib"
```

```

FEATURE_LIST_PATH = "feature_names.txt"

dump(best_overall_est, MODEL_PATH)
with open(FEATURE_LIST_PATH, "w") as f:
    for c in num_features:
        f.write(c + "\n")

print(f"\nSaved model to: {MODEL_PATH}")
print(f"Saved feature list to: {FEATURE_LIST_PATH}")

def predict_activities(new_data, model_path=MODEL_PATH, feature_list_p
    """
    new_data: pandas DataFrame (same feature space) OR path to CSV with
    Returns: numpy array of predicted activity labels.
    """
    from joblib import load
    if isinstance(new_data, str):
        new_df = pd.read_csv(new_data)
    else:
        new_df = new_data.copy()

    with open(feature_list_path, "r") as f:
        feats = [ln.strip() for ln in f if ln.strip()]

    missing = [c for c in feats if c not in new_df.columns]
    if missing:
        raise ValueError(f"Missing required columns: {missing[:10]}{'.'
    new_df = new_df.reindex(columns=feats)
    model = load(model_path)
    return model.predict(new_df)

```

Saved model to: best_har_model.joblib
 Saved feature list to: feature_names.txt

```

In [12]: # Bagging & Boosting add-ons (separate from your existing `models`)
from sklearn.ensemble import (
    BaggingClassifier, ExtraTreesClassifier, AdaBoostClassifier, HistG
)
from sklearn.tree import DecisionTreeClassifier

# A separate dict so you don't have to touch your original `models`
bag_boost_models = {
    # Bagging family
    "Bagging (DT)": passthrough_pipeline(
        BaggingClassifier(
            estimator=DecisionTreeClassifier(random_state=42),
            n_estimators=200, bootstrap=True, n_jobs=-1, random_state=
        )
    ),
    "ExtraTrees": passthrough_pipeline(
        ExtraTreesClassifier(n_estimators=400, n_jobs=-1, random_state
    ),

```

```

# Boosting family
"AdaBoost": passthrough_pipeline(
    AdaBoostClassifier(
        estimator=DecisionTreeClassifier(max_depth=2, random_state=n_estimators=300, learning_rate=0.5, random_state=42
    )
),
"HistGradientBoosting": passthrough_pipeline(
    HistGradientBoostingClassifier(
        max_iter=300, learning_rate=0.1, random_state=42
    )
),
}

```

```

In [13]: bb_results = []
bb_reports = {}

for name, pipe in bag_boost_models.items():
    pipe.fit(X_train, y_train)
    y_pred = pipe.predict(X_test)
    acc = accuracy_score(y_test, y_pred)
    f1m = f1_score(y_test, y_pred, average="macro")
    bb_results.append({"model": name, "accuracy": acc, "f1_macro": f1m})
    bb_reports[name] = classification_report(y_test, y_pred, digits=3)
    print(f"\n== {name} (baseline) ==")
    print(f"Accuracy: {acc:.4f} | f1_macro: {f1m:.4f}")

bb_baseline_df = pd.DataFrame(bb_results).sort_values(
    by=["accuracy", "f1_macro"], ascending=False
)
print("\n== Bagging/Boosting Baseline Leaderboard ==")
print(bb_baseline_df.to_string(index=False))

# If you already have results_df from earlier, show a merged view for
if 'results_df' in globals():
    merged = pd.concat([results_df, bb_baseline_df], ignore_index=True
                      .sort_values(by=["accuracy", "f1_macro"], ascending=False)
    print("\n== Combined Leaderboard (incl. new baselines) ==")
    print(merged.to_string(index=False))

```

```

== Bagging (DT) (baseline) ==
Accuracy: 0.9721 | f1_macro: 0.9722

== ExtraTrees (baseline) ==
Accuracy: 0.9932 | f1_macro: 0.9937

== AdaBoost (baseline) ==
Accuracy: 0.9864 | f1_macro: 0.9869

== HistGradientBoosting (baseline) ==
Accuracy: 0.9939 | f1_macro: 0.9941

== Bagging/Boosting Baseline Leaderboard ==
      model  accuracy  f1_macro
HistGradientBoosting  0.993882  0.994067
      ExtraTrees  0.993202  0.993722
      AdaBoost   0.986404  0.986879
      Bagging (DT)  0.972128  0.972211

== Combined Leaderboard (incl. new baselines) ==
      model  accuracy  f1_macro
      HistGradientBoosting  0.993882  0.994067
      ExtraTrees  0.993202  0.993722
      SVM (RBF) (tuned)  0.989803  0.990592
      SVM (RBF)   0.989803  0.990587
      RandomForest  0.987763  0.987936
      AdaBoost   0.986404  0.986879
      Logistic Regression  0.985724  0.986825
      Logistic Regression (tuned)  0.985724  0.986825
      RandomForest (tuned)  0.985044  0.985288
      RidgeClassifier (Adaline-like)  0.980965  0.982162
      Perceptron   0.980965  0.981771
      Bagging (DT)   0.972128  0.972211
      KNN (k=15) (tuned)  0.970768  0.972094
      KNN (k=15)   0.956492  0.957820
      DecisionTree  0.938817  0.937979

```

```

In [14]: from sklearn.model_selection import RandomizedSearchCV
         from scipy.stats import loguniform, randint as sp_randint

# Smaller n_iter to keep runtime reasonable; increase later if you want
QUICK = True
iters = (12 if QUICK else 30)

bb_search_specs = {
    "ExtraTrees": {
        "pipe": bag_boost_models["ExtraTrees"],
        "params": {
            "clf__n_estimators": sp_randint(200, 601),
            "clf__max_depth": [None, 10, 20, 30],
            "clf__max_features": ["sqrt", "log2", None],
            "clf__min_samples_split": sp_randint(2, 11),
            "clf__min_samples_leaf": sp_randint(1, 5),
        }
    }
}

```

```

        },
        "n_iter": iters
    },
    "Bagging (DT)": {
        "pipe": bag_boost_models["Bagging (DT)"],
        "params": {
            "clf_n_estimators": sp_randint(100, 401),
            "clf_bootstrap": [True, False],
            "clf_estimator_max_depth": [None, 5, 10, 20],
            "clf_estimator_min_samples_split": sp_randint(2, 11),
            "clf_estimator_min_samples_leaf": sp_randint(1, 5),
        },
        "n_iter": iters
    },
    "AdaBoost": {
        "pipe": bag_boost_models["AdaBoost"],
        "params": {
            "clf_n_estimators": sp_randint(100, 601),
            "clf_learning_rate": loguniform(1e-2, 1.0),
            "clf_estimator_max_depth": [1, 2, 3],
        },
        "n_iter": iters
    },
    "HistGradientBoosting": {
        "pipe": bag_boost_models["HistGradientBoosting"],
        "params": {
            "clf_learning_rate": loguniform(1e-2, 3e-1),
            "clf_max_iter": sp_randint(150, 401),
            "clf_max_leaf_nodes": sp_randint(15, 65),
            "clf_max_depth": [None, 6, 8, 10],
            "clf_min_samples_leaf": sp_randint(10, 60),
            "clf_l2_regularization": loguniform(1e-4, 1.0),
        },
        "n_iter": (iters + 6 if QUICK else 40)
    },
}
}

bb_tuned_models = {}
for name, spec in bb_search_specs.items():
    print(f"\n== RandomizedSearchCV for {name} ==")
    search = RandomizedSearchCV(
        estimator=spec["pipe"],
        param_distributions=spec["params"],
        n_iter=spec["n_iter"],
        cv=cv,
        scoring="f1_macro",
        random_state=42,
        n_jobs=-1,
        verbose=1,
        refit=True,
        error_score="raise"
    )

```

```

    search.fit(X_train, y_train)
    print(f"Best params for {name}: {search.best_params_}")
    print(f"CV best f1_macro: {search.best_score_:.4f}")
    bb_tuned_models[f"{name} (tuned)"] = search.best_estimator_

# Evaluate tuned ones
bb_more_results = [eval_and_log(nm, est) for nm, est in bb_tuned_models]
bb_more_df = pd.DataFrame(
    [{k: v for k, v in r.items() if k != "estimator"} for r in bb_more_results]
).sort_values(by=["accuracy", "f1_macro"], ascending=False)

print("\n== Bagging/Boosting Tuned Leaderboard ==")
print(bb_more_df.to_string(index=False))

# Optional: merge with previous overall leaderboard if available
if 'results_df' in globals():
    combined_df = pd.concat([results_df, bb_baseline_df, bb_more_df],
                           .sort_values(by=["accuracy", "f1_macro"], ascending=False))
    print("\n== Updated GRAND LEADERBOARD (incl. bagging/boosting) ==")
    print(combined_df.to_string(index=False))

```

```

== RandomizedSearchCV for ExtraTrees ==
Fitting 3 folds for each of 12 candidates, totalling 36 fits
Best params for ExtraTrees: {'clf__max_depth': 30, 'clf__max_features': 'sqrt', 'clf__min_samples_leaf': 1, 'clf__min_samples_split': 4, 'clf__n_estimators': 563}
CV best f1_macro: 0.9824

== RandomizedSearchCV for Bagging (DT) ==
Fitting 3 folds for each of 12 candidates, totalling 36 fits
Best params for Bagging (DT): {'clf__bootstrap': True, 'clf__estimator__max_depth': None, 'clf__estimator__min_samples_leaf': 3, 'clf__estimator__min_samples_split': 3, 'clf__n_estimators': 188}
CV best f1_macro: 0.9626

== RandomizedSearchCV for AdaBoost ==
Fitting 3 folds for each of 12 candidates, totalling 36 fits
Best params for AdaBoost: {'clf__estimator__max_depth': 3, 'clf__learning_rate': np.float64(0.39181943471417396), 'clf__n_estimators': 370}
CV best f1_macro: 0.9910

== RandomizedSearchCV for HistGradientBoosting ==
Fitting 3 folds for each of 18 candidates, totalling 54 fits
Best params for HistGradientBoosting: {'clf__l2_regularization': np.float64(0.12604664585649453), 'clf__learning_rate': np.float64(0.24420460844911424), 'clf__max_depth': 6, 'clf__max_iter': 377, 'clf__max_leaf_nodes': 28, 'clf__min_samples_leaf': 40}
CV best f1_macro: 0.9917

== ExtraTrees (tuned) ==
Accuracy: 0.9925 | f1_macro: 0.9931

== Bagging (DT) (tuned) ==

```

```

Accuracy: 0.9714 | f1_macro: 0.9712

== AdaBoost (tuned) ==
Accuracy: 0.9959 | f1_macro: 0.9961

== HistGradientBoosting (tuned) ==
Accuracy: 0.9925 | f1_macro: 0.9929

== Bagging/Boosting Tuned Leaderboard ==
      model accuracy f1_macro
AdaBoost (tuned) 0.995921 0.996052
ExtraTrees (tuned) 0.992522 0.993093
HistGradientBoosting (tuned) 0.992522 0.992915
Bagging (DT) (tuned) 0.971448 0.971197

== Updated GRAND LEADERBOARD (incl. bagging/boosting) ==
      model accuracy f1_macro
AdaBoost (tuned) 0.995921 0.996052
HistGradientBoosting 0.993882 0.994067
ExtraTrees 0.993202 0.993722
ExtraTrees (tuned) 0.992522 0.993093
HistGradientBoosting (tuned) 0.992522 0.992915
SVM (RBF) (tuned) 0.989803 0.990592
SVM (RBF) 0.989803 0.990587
RandomForest 0.987763 0.987936
AdaBoost 0.986404 0.986879
Logistic Regression 0.985724 0.986825
Logistic Regression (tuned) 0.985724 0.986825
RandomForest (tuned) 0.985044 0.985288
RidgeClassifier (Adaline-like) 0.980965 0.982162
Perceptron 0.980965 0.981771
Bagging (DT) 0.972128 0.972211
Bagging (DT) (tuned) 0.971448 0.971197
KNN (k=15) (tuned) 0.970768 0.972094
KNN (k=15) 0.956492 0.957820
DecisionTree 0.938817 0.937979

```

```

In [15]: # Pick best among the new/tuned models and show confusion
if 'bb_more_df' in globals() and not bb_more_df.empty:
    best_bb_name = bb_more_df.iloc[0]["model"]
    print(f"\n== Best Bagging/Boosting Model: {best_bb_name} ==")
    best_bb_est = {**bag_boost_models, **bb_tuned_models}[best_bb_name]
    from sklearn.metrics import ConfusionMatrixDisplay
    fig, ax = plt.subplots(figsize=(8,6))
    ConfusionMatrixDisplay.from_predictions(
        y_test, best_bb_est.predict(X_test),
        xticks_rotation=45, cmap="Blues", colorbar=False, ax=ax
    )
    ax.set_title(f"Confusion Matrix for {best_bb_name}")
    fig.tight_layout()
    plt.show()

```

```

== Best Bagging/Boosting Model: AdaBoost (tuned) ==

```

Confusion Matrix for AdaBoost (tuned)

